

▼ Ozone Pollution: EDA, Cleaning & Analysis

Purpose: Clean and validate EPA ozone monitoring data, compute daily maximum 8-hour ozone concentrations, analyze trends across regions and time, examine weekday/weekend effects, and produce a geospatial heatmap of high ozone concentrations.

```
import pandas as pd
ozone = pd.read_csv('/content/ozone.csv')
ozone.head()
```

	Date	Source	Site ID	POC	Daily Max 8-hour Ozone Concentration	Units	Daily AQI Value	Local Site Name	Daily Obs Count	Percent Complete	Method Code	Other
0	/2024	AQS	60010007	1	0.031	ppm	29.0	Livermore	17	100.0	47.0	418
1	01/02/2024	AQS	60010007	1	0.037	ppm	34.0	Livermore	17	100.0	47.0	418
2	/2024	AQS	60010007	1	NaN	ppm	30.0	Livermore	17	100.0	47.0	418
3	January 04/2024	AQS	60010007	1	0.026	ppm	24.0	Livermore	17	100.0	47.0	418
4	January 05/2024	AQS	60010007	1	0.027	ppm	25.0	Livermore	17	100.0	47.0	418

Next steps: [Generate code with ozone](#) [New interactive sheet](#)

```
ozone.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 29939 entries, 0 to 29938
Data columns (total 17 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   Date            29939 non-null   object 
 1   Source          29939 non-null   object 
 2   Site ID         29939 non-null   int64  
 3   POC             29939 non-null   int64  
 4   Daily Max 8-hour Ozone Concentration 28457 non-null   float64
 5   Units           29939 non-null   object 
 6   Daily AQI Value 28469 non-null   float64
 7   Local Site Name 29939 non-null   object 
 8   Daily Max 8-hour Ozone Concentration 29939 non-null   float64
 9   Daily AQI Value 29939 non-null   float64
 10  Local Site Name 29939 non-null   object 
 11  Daily Max 8-hour Ozone Concentration 29939 non-null   float64
 12  Daily AQI Value 29939 non-null   float64
 13  Local Site Name 29939 non-null   object 
 14  Daily Max 8-hour Ozone Concentration 29939 non-null   float64
 15  Daily AQI Value 29939 non-null   float64
 16  Local Site Name 29939 non-null   object 
 17  Daily Max 8-hour Ozone Concentration 29939 non-null   float64
 18  Daily AQI Value 29939 non-null   float64
 19  Local Site Name 29939 non-null   object 
```

```

8 Daily Obs Count           29939 non-null   int64
9 Percent Complete          29939 non-null   float64
10 Method Code               26406 non-null   float64
11 CBSA Code                 27935 non-null   float64
12 CBSA Name                  27935 non-null   object
13 County FIPS Code          29939 non-null   int64
14 County                      29939 non-null   object
15 Site Latitude              29938 non-null   float64
16 Site Longitude             29938 non-null   float64
dtypes: float64(7), int64(4), object(6)
memory usage: 3.9+ MB

```

```
# Count of missing values in each column
ozone.isna().sum()
```

	0
Date	0
Source	0
Site ID	0
POC	0
Daily Max 8-hour Ozone Concentration	1482
Units	0
Daily AQI Value	1470
Local Site Name	0
Daily Obs Count	0
Percent Complete	0
Method Code	3533
CBSA Code	2004
CBSA Name	2004
County FIPS Code	0
County	0
Site Latitude	1
Site Longitude	1

dtype: int64

```
import numpy as np
```

```
# Standardize column names
ozone.columns = ozone.columns.str.strip().str.lower().str.replace(' ', '_')
```

```
# Rename to consistent schema
ozone = ozone.rename(columns={
    'site_id': 'station_id',
    'local_site_name': 'site_name',
    'daily_max_8-hour_ozone_concentration': 'o3_ppb',
    'daily_aqi_value': 'aqi',
    'cbsa_name': 'region',
    'county': 'county',
    'site_latitude': 'latitude',
```

```
'site_longitude': 'longitude'
})
```

```
ozone.head()
```

0	/2024	AQS	60010007	1	0.031	ppm	29.0	Livermore	17	100.0	
1	01/02/2024	AQS	60010007	1	0.037	ppm	34.0	Livermore	17	100.0	
2	/2024	AQS	60010007	1	NaN	ppm	30.0	Livermore	17	100.0	
3	January 04/2024	AQS	60010007	1	0.026	ppm	24.0	Livermore	17	100.0	
4	January 05/2024	AQS	60010007	1	0.027	ppm	25.0	Livermore	17	100.0	

Next steps: [Generate code with ozone](#) [New interactive sheet](#)

```
# Parse date column
ozone['date'] = pd.to_datetime(ozone['date'], errors='coerce')
```

```
/tmp/ipython-input-2728457329.py:2: UserWarning: Could not infer format, so each element will be parsed
ozone['date'] = pd.to_datetime(ozone['date'], errors='coerce')
```

```
# Drop rows with missing ozone values
ozone = ozone.dropna(subset=['o3_ppb'])
```

```
# Check all units are consistent
print(ozone['units'].value_counts())

if ozone['units'].nunique() > 1:
    print("Multiple units detected - conversion needed!")
else:
    print("All units are consistent")
```

```
units
ppm    28457
Name: count, dtype: int64
All units are consistent
```

```
# Remove duplicates
```

```
ozone = ozone.drop_duplicates(subset=['station_id', 'date'])
```

```
# Identify missing metadata
missing_region = ozone['region'].isna().sum()
missing_coords = ozone[['latitude','longitude']].isna().sum()

print("Missing region entries:", missing_region)
print("Missing coordinates:", missing_coords.sum())
```

Missing region entries: 1589

Missing coordinates: 0

```
# Quick summary
print(ozone.describe())
```

	date	station_id	poc	\	
count	23745	2.374500e+04	23745.000000		
mean	2024-07-02 00:26:41.010738944	6.037171e+07	1.000547		
min	2024-01-01 00:00:00	6.001001e+07	1.000000		
25%	2024-04-02 00:00:00	6.019500e+07	1.000000		
50%	2024-07-03 00:00:00	6.037120e+07	1.000000		
75%	2024-10-01 00:00:00	6.059500e+07	1.000000		
max	2024-12-31 00:00:00	6.067001e+07	2.000000		
std	NaN	2.000220e+05	0.023392		
	o3_ppb	aqi	daily_obs_count	percent_complete	\
count	23745.000000	22535.000000	23745.000000	23745.000000	
mean	0.044550	46.869536	26.908697	99.402106	
min	0.001000	1.000000	1.000000	6.000000	
25%	0.035000	31.000000	17.000000	100.000000	
50%	0.042000	39.000000	17.000000	100.000000	
75%	0.053000	49.000000	17.000000	100.000000	
max	0.114000	209.000000	1000.000000	100.000000	
std	0.014874	26.569345	94.562316	3.519204	
	method_code	cbsa_code	county_fips_code	latitude	\
count	20900.000000	22156.000000	23745.000000	23745.000000	
mean	96.243636	31389.267016	37.010107	36.326261	
min	47.000000	12540.000000	1.000000	32.676180	
25%	87.000000	23420.000000	19.000000	34.066590	
50%	87.000000	31080.000000	37.000000	36.634225	
75%	87.000000	40900.000000	59.000000	38.201850	
max	199.000000	46380.000000	67.000000	40.776780	
std	41.479428	10557.967471	19.958428	2.124220	
	longitude				
count	23745.000000				
mean	-119.526684				
min	-124.179490				
25%	-121.338353				
50%	-119.658339				
75%	-118.127200				
max	-114.602090				
std	2.083342				

Key Analysis


Exploratory analysis and spatial/temporal insights:

1. Trend over time and regions

```
import matplotlib.pyplot as plt

ozone.groupby('date')['o3_ppb'].mean().plot(figsize=(10,5), title='Average Daily Max 8-hour Ozone in Ca')
plt.ylabel('Ozone (ppb)')
```

```
plt.show()
```



2. Regions with consistently high ozone

```
region_stats = (
    ozone.groupby('region')['o3_ppb']
    .agg(['mean', 'median', 'max', 'count'])
    .sort_values('mean', ascending=False)
)
print(region_stats.head(10))
```


region	mean	median	max	count
Riverside-San Bernardino-Ontario, CA	0.052377	0.049	0.113	3133
El Centro, CA	0.049987	0.050	0.114	985
Bakersfield, CA	0.049072	0.047	0.096	2212
Los Angeles-Long Beach-Anaheim, CA	0.047479	0.044	0.114	3934
Merced, CA	0.046677	0.045	0.084	285
Fresno, CA	0.045716	0.044	0.093	1978
Hanford-Corcoran, CA	0.045367	0.045	0.088	267
Madera, CA	0.045188	0.044	0.082	581
Bishop, CA	0.044873	0.044	0.076	797
Truckee-Grass Valley, CA	0.043021	0.041	0.072	288

3. Weekday vs weekend comparison

```
ozone['weekday'] = ozone['date'].dt.day_name()
ozone['is_weekend'] = ozone['weekday'].isin(['Saturday', 'Sunday'])

import seaborn as sns
sns.boxplot(x='is_weekend', y='o3_ppb', data=ozone)
```

```
<Axes: xlabel='is_weekend', ylabel='o3_ppb'>
```



4. Geospatial heatmap (daily averages per station)

```
import os
import folium
from folium.plugins import HeatMap

station_summary = ozone.groupby(['station_id', 'latitude', 'longitude'])['o3_ppb'].mean().reset_index()

m = folium.Map(location=[station_summary['latitude'].mean(), station_summary['longitude'].mean()], zoom=10)
HeatMap(station_summary[['latitude', 'longitude', 'o3_ppb']].values, radius=10).add_to(m)

#Ensure the output directory exist before saving
output_dir = 'outputs'
os.makedirs(output_dir, exist_ok=True)
m.save(f'{output_dir}/ozone_heatmap.html')
m
```

```
from google.colab import files  
files.download(f'{output_dir}/ozone_heatmap.html')
```

Environmental and Policy Recommendations

1. Target High-Risk Regions

Focus air-quality management and policy interventions in:

Riverside–San Bernardino–Ontario

Visalia–Porterville

El Centro

These areas exceed state and EPA's recommended thresholds more often and should be prioritized for mitigation.

2. Public awareness

Alert residents during high-ozone days and encourage reduced traffic or outdoor activity.

